

## Task-D: Collinear features and their effect on linear models

```
In [7]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import numpy as np
from sklearn.datasets import load_iris
from sklearn.linear_model import SGDClassifier
from sklearn.model_selection import GridSearchCV
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [8]: data = pd.read_csv('task_d.csv')
```

```
In [9]: data.head()
```

Out[9]:

	x	y	z	x*x	2*y	2*z+3*x*x	w	target
0	-0.581066	0.841837	-1.012978	-0.604025	0.841837	-0.665927	-0.536277	0
1	-0.894309	-0.207835	-1.012978	-0.883052	-0.207835	-0.917054	-0.522364	0
2	-1.207552	0.212034	-1.082312	-1.150918	0.212034	-1.166507	0.205738	0
3	-1.364174	0.002099	-0.943643	-1.280666	0.002099	-1.266540	-0.665720	0
4	-0.737687	1.051772	-1.012978	-0.744934	1.051772	-0.792746	-0.735054	0

```
In [10]: X = data.drop(['target'], axis=1).values
Y = data['target'].values
```

### Doing perturbation test to check the presence of collinearity

#### Task: 1 Logistic Regression

1. **Finding the Correlation between the features**
  - a. check the correlation between the features
  - b. plot heat map of correlation matrix using seaborn heatmap
2. **Finding the best model for the given data**
  - a. Train Logistic regression on data(X,Y) that we have created in the above cell
  - b. Find the best hyper parameter alpha with hyper parameter tuning using k-fold cross validation (grid search CV or random search CV make sure you choose the alpha in log space)
  - c. Create a new Logistic regression with the best alpha (search for how to get the best hyper parameter value), name the best model as 'best\_model'
3. **Getting the weights with the original data**
  - a. train the 'best\_model' with X, Y
  - b. Check the accuracy of the model 'best\_model\_accuracy'
  - c. Get the weights W using best\_model.coef\_
4. **Modifying original data**
  - a. Add a noise (order of  $10^{-2}$ ) to each element of X and get the new data set X' ( $X' = X + e$ )
  - b. Train the same 'best\_model' with data (X', Y)
  - c. Check the accuracy of the model 'best\_model\_accuracy\_edited'
  - d. Get the weights W' using best\_model.coef\_
5. **Checking deviations in metric and weights**
  - a. find the difference between 'best\_model\_accuracy\_edited' and 'best\_model\_accuracy'
  - b. find the absolute change between each value of W and W'  $\Rightarrow |(W - W')|$
  - c. print the top 4 features which have higher % change in weights compare to the other feature

## Task: 2 Linear SVM

1. Do the same steps (2, 3, 4, 5) we have done in the above task 1.

**Do write the observations based on the results you get from the deviations of weights in both Logistic Regression and linear SVM**

# TASK 1 - Logistic Regression

- a) Find the correlation between the features

```
In [11]: #standardization
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_sc = scaler.fit_transform(X,Y)
```

```
In [14]: model = SGDClassifier(loss='log')
model.fit(X_sc,Y)
```

```
Out[14]: SGDClassifier(loss='log')
```

```
In [13]: coef = model.coef_
for i,w in enumerate(coef[0]):
    print(f" weight {i}: {w}")
```

```
weight 0: 3.0371020053987197
weight 1: -2.5088767630268753
weight 2: 5.987294502945964
weight 3: 2.8943122925505875
weight 4: -2.5088767630268753
weight 5: 3.3277669089815243
weight 6: 4.977612635108219
```

**adding small error to each datapoint, which is taken from random variable with mean 0 and small variance**

```
In [15]: X_er = X_sc + np.random.normal(0,0.01,X_sc.shape)
```

```
In [16]: model.fit(X_er,Y)
coef_er = model.coef_
for i,w in enumerate(coef_er[0]):
    print(f" weight {i}: {w}")
```


```
weight 0: 11.021544146623631
weight 1: -12.896856095805347
weight 2: 21.094340010412246
weight 3: 10.775889569431406
weight 4: -12.894270547653974
weight 5: 12.199550371267506
weight 6: 9.785019743182849
```

**Before and after pertubation, the weights obtained are significantly different and therefore, it can be concluded that the features are collinear and cannot be used for feature importance**

b)seaborn heatmap of correlation matrix

```
In [39]: figure = plt.subplots(figsize=(15,7))
sns.heatmap(data.corr(), annot=True, cmap='BuGn', linewidths=3, linecolor='black',
```

```
Out[39]: <matplotlib.axes._subplots.AxesSubplot at 0x229de4e36a0>
```



**looking at the correlation matrix, following features have strong linear relationship:**

x,y

z, x\*\*

w,z

z, 2z+x\*x

w,x\*x (moderate)

w,2z+x\*x(moderate)

2.Find best model for given data

a. Train Logistic regression on data(X,Y) that we have created in the above cell

```
In [17]: clf = SGDClassifier(loss='log')
clf.fit(X,Y)
```

```
Out[17]: SGDClassifier(loss='log')
```

b. Find the best hyper parameter alpha with hyper parameter tuning using k-fold cross validation (grid search CV or random search CV make sure you choose the alpha in log space)

```
In [18]: alpha = np.logspace(0,2)
hyperparameter = [{'alpha': alpha}]
grid_search = GridSearchCV(clf,hyperparameter,cv=5)
```

c. Create a new Logistic regression with the best alpha (search for how to get the best hyper parameter value), name the best model as 'best\_model'

```
In [19]: best_model = grid_search.fit(X,Y)
```

```
In [20]: #best hyperparameter value
best_model.best_params_
```

```
Out[20]: {'alpha': 1.0}
```

## The best hyperparameter value was found to be 1.0 through grid search - 5 fold cv

3. Getting the weights with the original data

a. train the 'best\_model' with X, Y

```
In [21]: best_model = SGDClassifier(loss='log',alpha=1.0)
best_model.fit(X,Y)
```

```
Out[21]: SGDClassifier(alpha=1.0, loss='log')
```

b. Check the accuracy of the model 'best\_model\_accuracy'

```
In [24]: best_model_accuracy = best_model.score(X,Y)
best_model_accuracy
```

```
Out[24]: 0.99
```

## After training the model with best hyperparameter value, a very high accuracy (0.99) is obtained

c. Get the weights W using best\_model.coef\_

```
In [26]: W = best_model.coef_
```

#### 4. Modifying original data

a. Add a noise (order of  $10^{-2}$ ) to each element of  $X$  and get the new data set  $X'$  ( $X' = X + e$ )

```
In [27]: X_err = X + (np.random.randint(0,100) * (10 ** -2))
```

b. Train the same 'best\_model' with data ( $X'$ ,  $Y$ )

```
In [28]: best_model.fit(X_err,Y)
```

```
Out[28]: SGDClassifier(alpha=1.0, loss='log')
```

c. Check the accuracy of the model 'best\_model\_accuracy\_edited'

```
In [31]: best_model_accuracy_edited = best_model.score(X_err,Y)
best_model_accuracy_edited
```

```
Out[31]: 0.99
```

## accuracy of model remains same even after addition of outlier.

d. Get the weights  $W'$  using `best_model.coef_`

```
In [32]: W_ = best_model.coef_
```

#### 5. Checking deviations in metric and weights

a. find the difference between 'best\_model\_accuracy\_edited' and 'best\_model\_accuracy'

```
In [33]: best_model_accuracy_edited - best_model_accuracy
```

```
Out[33]: 0.0
```

b. find the absolute change between each value of  $W$  and  $W'$   $\Rightarrow |(W-W')|$

```
In [34]: W - W_
```

```
Out[34]: array([[0.00427724, 0.01273821, 0.01089136, 0.00373512, 0.01273821,
                0.00431907, 0.01095538]])
```

c. print the top 4 features which have higher % change in weights compare to the other feature

```
In [255]: h = W - W_  
h = np.sort(h)  
h = h.reshape(-1,1)  
h = h[::-1]
```

```
In [256]: h[0:4]
```

```
Out[256]: array([[0.03122248],  
                [0.03122248],  
                [0.02837311],  
                [0.02643868]])
```

## TASK 2

2. Finding the best model for the given data

a. Train linear SVM on data(X,Y) that we have created in the above cell

```
In [35]: clf = SGDClassifier(loss='hinge')  
clf.fit(X,Y)
```

```
Out[35]: SGDClassifier()
```

b. Find the best hyper parameter alpha with hyper parameter tuning using k-fold cross validation (grid search CV or random search CV make sure you choose the alpha in log space)

```
In [36]: alpha = np.logspace(0,2)  
hyperparameter = [{'alpha': alpha}]  
grid_search = GridSearchCV(clf,hyperparameter,cv=5)
```

c. Create a new Linear SVM with the best alpha (search for how to get the best hyper parameter value), name the best model as 'best\_model'

```
In [37]: best_model = grid_search.fit(X,Y)
```

```
In [38]: best_model.best_params_
```

```
Out[38]: {'alpha': 1.0}
```

3. Getting the weights with the original data

a. train the 'best\_model' with X, Y

```
In [44]: best_model = SGDClassifier(loss='hinge',alpha=1.0)
best_model.fit(X,Y)
```

```
Out[44]: SGDClassifier(alpha=1.0)
```

b. Check the accuracy of the model 'best\_model\_accuracy'

```
In [45]: best_model_accuracy = best_model.score(X,Y)
best_model_accuracy
```

```
Out[45]: 1.0
```

## 100% accuracy is obtained with best hyperparameter value - 1.0

c. Get the weights W using best\_model.coef\_

```
In [46]: W = best_model.coef_
```

4.Modifying original data

a. Add a noise(order of  $10^{-2}$ ) to each element of X and get the new data set X' ( $X' = X + e$ )

```
In [47]: X_err = X + (np.random.randint(0,100) * (10 ** -2))
```

b. Train the same 'best\_model' with data (X', Y)

```
In [48]: best_model.fit(X_err,Y)
```

```
Out[48]: SGDClassifier(alpha=1.0)
```

c. Check the accuracy of the model 'best\_model\_accuracy\_edited'

```
In [51]: best_model_accuracy_edited = best_model.score(X_err,Y)
```

```
Out[51]: 0.0
```

## Accuracy of the model is not impacted by outliers

d. Get the weights W' using best\_model.coef\_

```
In [52]: W_ = best_model.coef_
```

5.Checking deviations in metric and weights



a. find the difference between 'best\_model\_accuracy\_edited' and 'best\_model\_accuracy'

```
In [53]: best_model_accuracy - best_model_accuracy_edited
```

```
Out[53]: 0.0
```

b. find the absolute change between each value of  $W$  and  $W'$  ==>  $|W-W'|$

```
In [54]: W - W_
```

```
Out[54]: array([[ -0.00237432, -0.00487535, -0.00322463, -0.0040665 , -0.00487535,
                -0.00372149,  0.00720876]])
```

c. print the top 4 features which have higher % change in weights compare to the other feature

```
In [55]: h = W - W_
h = np.sort(h)
h = h.reshape(-1,1)
h = h[::-1]
```

```
In [56]: h[0:4]
```

```
Out[56]: array([[ 0.00720876],
                [-0.00237432],
                [-0.00322463],
                [-0.00372149]])
```

**Here the change in value of weights is in the range of  $10^{-2}$**